

# Guaranteed Recovery of Sufficiently Unambiguous Clusters

Kayvon Mazooji\*

Ilan Shomorony†

## Abstract

Clustering is often a challenging problem because of the inherent ambiguity in what the “correct” clustering should be. Even when the number of clusters  $K$  is known, this ambiguity can still exist, particularly when there is variation in density among different clusters, and clusters have multiple relatively separated regions of high density. In this paper, we introduce a new algorithm that provably recovers a  $K$ -clustering  $C$  if it satisfies a relaxed cluster separability condition. This natural cluster separability condition formalizes the situation when two high density regions within a cluster are separable enough that they look more like two distinct clusters than two truly distinct clusters in  $C$ . The algorithm first identifies  $K$  partial clusters (or “seeds”) using a density-based approach, and then adds unclustered points to the initial  $K$  partial clusters in a greedy manner to form a complete clustering. This framework allows for the guaranteed recovery of sufficiently unambiguous non-convex clusters with arbitrarily many relatively separated regions of high density, and arbitrary variation in density among different clusters. The algorithm requires little parameter selection, and displays improved performance on many datasets compared to state of the art algorithms for non-convex cluster recovery.

## 1 Introduction

More often than not, there is significant ambiguity in what the “correct” clustering is for a dataset. Even in the case when the number of clusters  $K$  is known, this ambiguity often still exists because a cluster can have multiple relatively separated regions of high density, and clusters can have very different densities. These difficulties cause many algorithms to incorrectly separate a true cluster into multiple clusters, while merging true clusters or failing to detect sufficiently prominent true clusters of low density. These issues are compounded by the well-known fact that many clustering algorithms are ineffective at identifying non-convex clusters, which are present in many applications including image segmentation [1], geospatial data [2, 3], and time series data [4].

The most widely used paradigms for finding non-convex clusters are spectral clustering, and density-

based clustering. Spectral clustering algorithms perform dimensionality reduction on the data points, and run a simple clustering algorithm such as  $K$ -means to cluster the low dimensional data [5]. Density-based clustering algorithms find regions of high point density, and then output a set of high density clusters according to some criterion [6–9]. While spectral clustering is very well studied from a theoretical standpoint, most widely used density-based algorithms do not come with theoretical guarantees on when a clustering is recoverable. In this paper, we try to overcome the difficulties in  $K$ -clustering described in the previous paragraph from a theoretical and practical standpoint, through the lens of density-based clustering.

In specific, we design an efficient algorithm that provably recovers a  $K$ -clustering  $C$  under a novel and intuitive separability condition. This separability condition on  $C$  characterizes the situation when two high density regions within the same cluster in  $C$  look more like two distinct cluster than two truly distinct clusters in  $C$ . This yields the provable reconstruction of sufficiently unambiguous  $K$ -clusterings that can have

- clusters with arbitrarily many relatively separated regions of high density
- arbitrary variation in density among different clusters
- arbitrary variation in density within clusters
- arbitrarily shaped clusters.

The algorithm for accomplishing this is somewhat complex, but has a computationally efficient runtime. It works by first finding a small subset of points called a “seed” from each cluster, and then expanding the seeds to form clusters. The algorithm requires little parameter selection, and delivers improved performance on artificial and benchmark datasets compared to state of the art algorithms for non-convex cluster recovery.

We begin by discussing related work and introducing background information and notation. We then introduce several novel cluster separability conditions of varying degrees of strictness, including our primary separability condition. Next, we present our algorithm, and prove that it recovers the true clustering under each of these separability conditions. Finally, we compare the performance of our algorithm with spectral clustering,  $K$ -means, HDBSCAN, and OPTICS on artificial

---

\*University of Illinois Urbana-Champaign.

†University of Illinois Urbana-Champaign.

datasets and real-world benchmark datasets.

## 2 Related Work

The first widely-used density-based clustering algorithm was DBSCAN [6], which was introduced in 1996. Since then, many algorithms have been designed to improve upon various aspects of DBSCAN [7–14]. In addition to DBSCAN, the other most widely used density-based clustering algorithms are OPTICS [7] and HDBSCAN [8], which were both designed to improve upon DBSCAN’s ability to output clusters of varying density. On the theoretical side, there has been a recent line of work studying  $\lambda$ -density level set estimation using DBSCAN [14–16].

Our algorithm works by first identifying  $K$  partial clusters (or “seeds”) using a density-based approach, and then adding unclustered points to the initial  $K$  partial clusters in a greedy manner to form a complete clustering. To extract these seeds, we sequentially find seeds of decreasing density. The intuitive idea of sequentially finding disjoint clusters of decreasing density has been employed in various work [11, 12]. However, to the best of our knowledge, our algorithm for finding  $K$  seeds and our algorithm for expanding the seeds to form complete clusters have not previously appeared in the literature. Furthermore, we are not aware of any existing mathematical analysis of density-based clustering algorithms that give guarantees similar to those presented in this paper.

## 3 Preliminaries

We use  $X$  to denote a set of data points. The distance between points  $x$  and  $y$  is denoted by  $d(x, y)$ . As in many density-based clustering papers, the measure of density at a point  $x \in X$  is determined by an integer  $N_p$ , and is defined as  $1/\epsilon_{N_p}(x)$  where  $\epsilon_{N_p}(x)$  is the minimum distance  $\epsilon$  such that there are  $N_p$  points in  $X$  at distance at most  $\epsilon$  from  $x$  (including  $x$  itself). In other words,  $\epsilon_{N_p}(x)$  is the distance from  $x$  to its  $(N_p - 1)$ th nearest neighbor in  $X$ . We call  $\epsilon_{N_p}(x)$  the sparsity at  $x$ .

A cluster is simply a set of points in  $X$ . A clustering  $C$  of  $X$  is a set of disjoint clusters, where every point in  $X$  belongs to exactly one cluster (i.e. a partitioning of  $X$  into clusters). A  $K$ -clustering is a clustering with  $K$  clusters. A partial clustering of  $X$  is a set of disjoint clusters whose union does not necessarily include all points in  $X$ . We say that a clustering  $C$  extends (or is an extension of) a partial clustering  $C'$  if there exists a bijective function  $f$  that maps the clusters in  $C'$  to the clusters in  $C$  such that for each cluster  $c' \in C'$ ,  $c'$  is a non-empty subset of the cluster  $f(c') \in C$ .

We say a point  $x_1$  is  $\epsilon$ -connected to a point  $x_t$  if there exists some sequence of points  $x_1, x_2, \dots, x_t$  such

that  $x_{i+1}$  is distance at most  $\epsilon$  from  $x_i$  for  $1 < i < t - 1$  and  $\epsilon_{N_p}(x_i) \leq \epsilon$  for  $1 < i < t$ . A set of points  $c$  is called  $\epsilon$ -connected if every pair of points in  $c$  is  $\epsilon$ -connected.

For a given  $\epsilon$ , a set of points  $c$  is called an  $\epsilon$ -cluster if it is a maximal  $\epsilon$ -connected set of points. In other words, any point that is  $\epsilon$ -connected to a point in  $c$  is included in  $c$ . A set of points  $c$  is called a maximal cluster if it is an  $\epsilon$ -cluster for some  $\epsilon$ . For a given  $\epsilon$ , it is helpful consider the graph that is formed where each node corresponds to a data point, and an edge is drawn between two nodes if the corresponding points  $x_1, x_2$  are such that  $\epsilon_{N_p}(x_1) \leq \epsilon$ ,  $\epsilon_{N_p}(x_2) \leq \epsilon$ , and  $d(x_1, x_2) \leq \epsilon$ . A set of points is an  $\epsilon$ -cluster if and only if it corresponds to a connected component in this graph.

The  $\epsilon$ -cluster centered at a point  $x$  is defined as the set of points that are  $\epsilon$ -connected to  $x$ , and is denoted by  $c^*(x, \epsilon)$ . If  $\epsilon < \epsilon_{N_p}(x)$ , then  $c^*(x, \epsilon) = \{x\}$ . The sparsity of a set of points  $c$  is defined as the minimum  $\epsilon$  such that  $c$  is  $\epsilon$ -connected, and is denoted by  $\epsilon^*(c)$ . The  $\epsilon$ -distance between points  $x$  and  $y$  is defined as the minimum  $\epsilon$  such that  $x$  is  $\epsilon$ -connected to  $y$ , and is denoted by  $\epsilon(x, y)$ . The minimum  $\epsilon$ -distance from a point  $x$  to a cluster  $c$  is defined as  $\epsilon(x, c) = \min_{y \in c} \epsilon(x, y)$ . The minimum  $\epsilon$ -distance from cluster  $c_1$  to cluster  $c_2$  is defined as  $\epsilon(c_1, c_2) = \min_{x \in c_1, y \in c_2} \epsilon(x, y)$ .

For a dataset  $X$  and density parameter  $N_p$ , the dendrogram  $G = (V, E)$  is a tree structure that gives all  $\epsilon$ -clusters in  $X$  for each  $\epsilon \geq 0$ .  $V$  and  $E$  are the sets of nodes and edges in  $G$  respectively. Each node  $v \in V$  corresponds to an  $\epsilon$ -cluster for some  $\epsilon$ , and the clusters corresponding to the children of  $v$  form the smallest possible partition of the maximal cluster corresponding to  $v$  into maximal clusters. Each value of  $\epsilon$  specifies a clustering given by all nodes in  $V$  corresponding to  $\epsilon$ -clusters. For a given  $x$  and  $N_p$ , the dendrogram is unique, and yields a hierarchy of possible clusterings, making it the foundational structure in hierarchical density-based clustering (HDBSCAN) [8].

For a node  $v \in V$ ,  $c(v)$  denotes the cluster corresponding to  $v$ . Each node  $v$  has a real number  $\epsilon(v)$  associated with it where  $\epsilon(v)$  is the smallest  $\epsilon$  such that  $c(v)$  is a  $\epsilon$ -cluster. For a leaf node  $v$  corresponding to the cluster  $\{x\}$ ,  $\epsilon(v) = \epsilon_{N_p}(x)$ . A node  $v_0$  has children  $v_1, v_2, \dots, v_i$  if  $c(v_1), c(v_2), \dots, c(v_i)$  form the smallest possible partition of  $c(v_0)$  composed of maximal clusters, which is guaranteed to be unique. The root node  $v_r$  of the resulting tree is such that  $c(v_r) = X$ .

Any  $S \subset V$  such that no node in  $S$  is a descendant of another node in  $S$  induces a (partial) clustering of  $X$  given by  $\{c(v) \text{ for } v \in S\}$ . Any such clustering is called a dendrogram clustering. In fact, any (partial) clustering that consists of maximal clusters is a dendrogram (partial) clustering since every maximal cluster corre-

sponds to a node in the dendrogram. The  $\epsilon$ -cut of a dendrogram is the (partial) clustering that includes all clusters corresponding to nodes  $v$  such that  $\epsilon(v) \leq \epsilon$ , and  $v$  has no parent  $v'$  with  $\epsilon(v') \leq \epsilon$ . For any integer  $K$ , there exists at most one partial clustering given by an  $\epsilon$ -cut of  $G$  that contains  $K$  clusters. The (partial) clustering output by DBSCAN for a given  $N_p$  and  $\epsilon$  can be formed by taking the  $\epsilon$ -cut of  $G$ , and adding any unclustered point  $x$  to the cluster  $c$  that minimizes  $d(x, c)$  if  $d(x, c) \leq \epsilon$  where  $d(x, c) = \min_{y \in c} d(x, y)$ .

## 4 Results

Recovering a clustering  $C$  of  $X$  with a known number of clusters  $K$  is often difficult because two high density regions within a cluster in  $C$  can look more like two distinct clusters than two truly distinct clusters in  $C$ . We introduce novel a cluster separability condition that does not allow this to happen, and design an algorithm that is guaranteed to recover  $C$  when the condition is satisfied. We begin this discussion by comparing and contrasting several notions of cluster separability.

**4.1 Weak Separability** The simplest density-based notion of separability is what we call weak separability.

**DEFINITION 4.1.** *For a given  $N_p$ ,  $C$  is called weakly separable if for each  $c \in C$ ,  $c$  is  $\epsilon$ -connected where  $\epsilon < \min_{c' \in C, c' \neq c} \epsilon(c, c')$ .*

Interestingly, a dataset  $X$  often has more than one weakly separable  $K$ -clustering. For example, suppose that  $N_p = 2$  and  $C = [\{1, 3, 5, 7.02, 9.02, 11.02\}, \{17, 18, 19, 20\}, \{22.01, 23.01, 24.01, 25.01\}]$ . Clearly,  $C$  is weakly separable. However,  $C' = [\{1, 3, 5\}, \{7.02, 9.02, 11.02\}, \{17, 18, 19, 20, 22.01, 23.01, 24.01, 25.01\}]$  is also a weakly separable 3-clustering. Intuitively,  $C$  is the correct clustering because the spacing between points is nearly uniform within each cluster in  $C$ , while there is a large relative gap in the middle of the third cluster in  $C'$ . There are many such cases where the intuitively correct clustering is weakly separable, but is not the unique weakly separable clustering. In fact, we prove that any set of maximal clusters that partition  $X$  (a dendrogram clustering) is weakly separable.

**COROLLARY 4.1.** *For a given  $N_p$ ,  $X$  has a unique weakly-separable  $K$ -clustering if and only if exactly one dendrogram  $K$ -clustering exists.*

*Proof.* This follows from Lemma 4.1.  $\square$

**LEMMA 4.1.** *For a given  $N_p$ ,  $C$  is weakly separable if and only if it is a dendrogram clustering.*

*Proof.* Suppose  $C$  is a dendrogram clustering, but is not weakly separable. Let  $G = (V, E)$  be the dendrogram. Then there exists some  $v, v' \in V$  such that  $c(v), c(v') \in C$ , we have that  $\epsilon(v) \geq \epsilon(c(v), c(v'))$ , then  $c(v)$  would include points from  $c(v')$  since  $c(v)$  is maximal. This contradicts the definition of a clustering.

If  $C$  is weakly separable, then each cluster in  $C$  is maximal. This follows because if a cluster  $c$  is not maximal, there exists a point  $x \notin c$  that is  $\epsilon$ -connected to  $c$  such that  $\epsilon < \epsilon^*(c)$ , which contradicts weak separability. Every maximal cluster has a corresponding node in  $G$ , so  $C$  is a dendrogram clustering.  $\square$

If there are two weakly separable clusterings for  $X$ , it is impossible to guarantee that we recover one of them if our only criteria is to find a weakly separable clustering. Because there usually does not exist a unique weakly separable clustering, weak separability alone as a sufficient condition for clustering recoverability is not general enough.

**4.2 Strong Separability** A similar, yet stronger notion of cluster separability is what we call strong separability. For a cluster  $c \in C$ , let  $X_c^*$  denote the set of all points  $x$  in  $c$  such that  $\epsilon_{N_p}(x) = \min_{y \in c} \epsilon_{N_p}(y)$ . In other words,  $X_c^*$  denotes the set of all points  $x$  in  $c$  that have highest density among points in  $c$ .

**DEFINITION 4.2.** *For a given  $N_p$ ,  $C$  is called strongly separable if there exists some  $A \in \mathbb{R}$  such that one of the following equivalent conditions holds.*

1. *For each  $c \in C$ ,  $c$  is  $(A \cdot \min_{x \in c} \epsilon_{N_p}(x))$ -connected, and  $A \cdot \min_{x \in c} \epsilon_{N_p}(x) < \min_{c' \in C, c' \neq c} \epsilon(c, c')$ .*
2. *For each  $c \in C$ ,  $c^*(x, A \cdot \epsilon_{N_p}(x)) = c$  for any  $x \in X_c^*$*

*The minimum such  $A$  is denoted by  $A^*(C)$ .*

The definition specifies for each cluster  $c$ , an  $\epsilon$  relative to the sparsity of the maximum density point in  $c$ , such that no cluster can be  $\epsilon$ -connected to  $c$ , yet all points in  $c$  must be  $\epsilon$ -connected. The condition is therefore naturally satisfied by clusterings that have arbitrarily shaped clusters with arbitrarily many relatively separated regions of high density and arbitrary variation in density among different clusters, so long as the clusters are separated enough relative to the sparsity values of their respective maximum density points. Unlike weak separability, if  $C$  is strongly separable for  $N_p$ , then it is the unique strongly separable clustering for  $N_p$  and can be found efficiently by Theorem 4.1, which follows immediately from Theorem 4.2 and Lemmas 4.4 and 4.3.

**THEOREM 4.1.** *If  $C$  is strongly separable for a given  $N_p$ , then  $C$  is the unique strongly separable clustering for  $N_p$ , and can be found in  $O(|X|^3 \log(|X|))$  time.*

Perhaps a more intuitive condition that implies strong separability is given in the following lemma, which for each cluster  $c \in C$ , bounds the maximum variation in density among points in  $C$ . Let  $\alpha(c) = \max_{x \in c} \epsilon_{N_p}(x) / \min_{y \in c} \epsilon_{N_p}(y)$  be the ratio of the sparsity of the minimum density point in  $X$  to that of the maximum density point in  $X$ .

**LEMMA 4.2.** *For a given  $N_p$ , if there exists some  $A \in \mathbb{R}$  such that for every  $c \in C$ ,  $\alpha(c) \leq A$ ,  $c$  is  $(\max_{x \in c} \epsilon_{N_p}(x))$ -connected, and  $A \cdot \min_{x \in c} \epsilon_{N_p}(x) < \min_{c' \in C, c' \neq c} \epsilon(c, c')$ , then  $C$  is strongly separable.*

*Proof.* The fact that for each  $c \in C$ ,  $\alpha(c) \leq A$  and  $c$  is  $(\max_{x \in c} \epsilon_{N_p}(x))$ -connected implies that for each  $c \in C$ ,  $c$  is  $(A \cdot \min_{x \in c} \epsilon_{N_p}(x))$ -connected.  $\square$

If  $C$  is strongly separable, then it is weakly separable as proved in Lemma 4.3. However, strong separability of  $C$  does not imply that  $C$  is the unique weakly separable  $|C|$ -clustering for  $N_p$ . Furthermore, there exist weakly separable clusterings that are not strongly separable. For example, for  $N_p = 2$ ,  $C = [\{1, 3, 5\}, \{8, 10, 11, 13\}]$  is weakly separable but not strongly separable since the points in the second cluster imply that  $A^*(C) \geq 2$ , but  $c^*(3, 2 \cdot \epsilon_{N_p}(3)) = c^*(3, 4) = X$ . Intuitively,  $C$  is the correct 2-clustering, thus showing that strong separability as a sufficient condition for clustering recovery is still not general enough. Therefore, we proceed to introduce a more inclusive sufficient condition for recoverability in the next subsection. Figure 1 gives a larger example of a weakly separable clustering that is not strongly separable.

**LEMMA 4.3.** *For a given  $N_p$ , if  $C$  is strongly separable, then it is weakly separable.*

*Proof.* This follows immediately from the first equivalent definition of strong separability.  $\square$

For a given  $N_p$ , it is in general not possible to recover the strongly separable clustering by simply finding an extension of the (partial) clustering given by the  $\epsilon$ -cut of the dendrogram that contains  $K$  clusters (if there exists such an  $\epsilon$ ) as proved in Lemma 4.1. Since DBSCAN follows this approach, it is not sufficient for finding the strongly separable clustering.

**LEMMA 4.1.** *There exists a strongly separable clustering for some  $N_p$  that does not extend any (partial) clustering given by an  $\epsilon$ -cut of the dendrogram.*

*Proof.* Recall that there can only be one  $\epsilon$ -cut of the dendrogram that gives a (partial)  $K$ -clustering. Let  $N_p = 3$ , and suppose that

$C = [\{1, 3, 5, 7.02, 9.02, 11.02\}, \{17, 18, 19, 20\}, \{22.01, 23.01, 24.01, 25.01\}]$ . This is a strongly separable clustering and  $A^*(C) = 2$ . The only  $\epsilon$ -cut that gives a (partial) clustering with three clusters is chosen by setting  $\epsilon = 2.01$ , and is given by  $C' = [\{3\}, \{9.02\}, \{17, 18, 19, 20, 22.01, 23.01, 24.01, 25.01\}]$ . The only weakly separable clustering that is an extension of  $C'$  is  $[\{1, 3, 5\}, \{7.02, 9.02, 11.02\}, \{17, 18, 19, 20, 22.01, 23.01, 24.01, 25.01\}]$ .  $\square$

In the clustering  $C$  used to prove Lemma 4.1, the distance separating the points 1, 3, 5 from the points 7.02, 9.02, 11.02 within the first cluster is 2.02, which is larger than the distance of 2.01 separating the clusters  $\{17, 18, 19, 20\}$  and  $\{22.01, 23.01, 24.01, 25.01\}$ . However, 2.02 is very similar to the distance of 2 separating the other pairs of adjacent points in the first cluster, while 2.01 is very large compared to the distance of 1 separating the adjacent points in the second cluster.  $C$  is therefore a more natural clustering than  $C'$  in a sense, but an  $\epsilon$ -cut is unable to capture  $C$  because it only considers absolute distances when separating clusters, without taking cluster density into account.

**4.3 Local Maximum Separability** The main conceptual contribution of this paper is a novel cluster separability condition we call local maximum separability (or LM-separability), along with a new algorithm that provably recovers any weakly separable, local maximum separable clustering. We call a point  $x \in X$  a local maximum if  $\epsilon_{N_p}(y) \geq \epsilon_{N_p}(x)$  for  $y \in X$  such that  $d(x, y) \leq \epsilon_{N_p}(x)$ . For a given density parameter  $N_p$ , the relative separability of a point  $x \in X$  to a point  $y \in X$  is the value of  $A$  such that  $A \cdot \epsilon_{N_p}(x) = \epsilon(x, y)$ , and is denoted by  $A(x, y)$ . Similarly, the relative separability of a point  $x \in X$  to a cluster  $c$  is given by  $\min_{y \in c} A(x, y)$ , and is denoted by  $A(x, c)$ .

**DEFINITION 4.3.** *For a given  $N_p$  and  $C$ , we use  $A^\ell(C)$  to denote the minimum  $A \in \mathbb{R}$  such that*

$$\min_{y \in c(x): \epsilon_{N_p}(y) < \epsilon_{N_p}(x)} A(x, y) \leq A,$$

*for every local maximum  $x \in X$  where there exists a  $y \in c(x)$  such that  $\epsilon_{N_p}(y) < \epsilon_{N_p}(x)$ , and*

$$\min_{y \in c(x): \epsilon_{N_p}(y) = \epsilon_{N_p}(x)} A(x, y) \leq A$$

*for every local maximum  $x \in X$  where  $\epsilon_{N_p}(x) = \min_{y \in c(x)} \epsilon_{N_p}(y)$  and there exists a  $y \in c(x)$  such that  $y \neq x$  and  $\epsilon_{N_p}(y) = \epsilon_{N_p}(x)$ .  $C$  is called local maximum separable (LM-separable) if*

$$A^\ell(C) < \min_{c, c' \in C} \min_{z \in X_c^*} A(z, c').$$

In specific, if  $C$  has only one local maximum per cluster, then  $C$  is trivially LM-separable.

LM-separability specifies that for every local maximum  $x \in X$  where there exists a  $y \in c(x)$  whose density is higher than that of  $x$ , the minimum relative separability of  $x$  to such a  $y$  is smaller than the relative separability of the highest density point in any cluster to another cluster. This is a natural notion of separability because if there exists a local maximum  $x$  where the minimum relative separability to such a point  $y \in c(x)$  is at least as high as the relative separability of a highest density point  $z \in X_c^*$  for some  $c \in C$  to another  $c' \in C$  (and therefore to a local maximum in  $c'$ ), then in a sense,  $x$  looks more like it belongs to a separate cluster from  $y$  than  $z$  looks like it belongs to a separate cluster from  $c'$ , and the correct  $K$ -clustering of  $X$  is ambiguous. LM-separability is more natural than strong separability because LM-separability and weak separability hold precisely when a clustering is unambiguous. In fact, strong separability implies LM-separability.

**LEMMA 4.4.** *For a given  $N_p$ , if  $C$  is strongly separable, then it is LM-separable.*

*Proof.* If there is only one local maximum per cluster, then this holds trivially. Suppose  $C$  is strongly separable but there exists a local maximum  $x$  such that  $\min_{y \in c(x): \epsilon_{N_p}(y) < \epsilon_{N_p}(x)} A(x, y) \geq \min_{c, c' \in C} \min_{z \in X_c^*} A(z, c')$ . Since  $\epsilon_{N_p}(x) > \epsilon_{N_p}(w)$  for all  $w \in X_{c(x)}^*$ , this implies that  $A(w, x) \geq \min_{c, c' \in C} \min_{z \in X_c^*} A(z, c')$  for all  $w \in X_{c(x)}^*$ , which in turn implies that  $A^*(C) \geq \min_{c, c' \in C} \min_{z \in X_c^*} A(z, c')$ . Thus, for the  $c, c', z$  that minimize  $\min_{c, c' \in C} \min_{z \in X_c^*} A(z, c')$ , we have that  $c^*(z, A^*(C) \cdot \epsilon_{N_p}(z))$  contains at least one point from  $c'$ . This is a contradiction to strong separability.  $\square$

Note that LM-separability does not imply weak separability. Let  $N_p = 2$ , and consider the clustering  $C = [\{7, 8, 10, 13, 21\}, \{17, 25, 27\}]$ .  $C$  is clearly not weakly separable because the clusters overlap, but is LM-separable, as 7, 8, 25, 27 are the only local maximums. Furthermore, weak separability together with LM-separability does not imply strong separability. Let  $N_p = 2$ , and consider the clustering  $C = [\{7, 8, 10, 13\}, \{17, 19, 21\}]$ .  $C$  is weakly separable, and is LM-separable as 7, 8, 17, 19, 21 are the only local maximums.  $C$  is not strongly separable because the first cluster implies that  $A^*(C) \geq 3$ , but  $c^*(19, 3 \cdot \epsilon_{N_p}(19)) = c^*(19, 6) = X$ . Figure 1 shows a larger example of a clustering that is not strongly separable, but is weakly separable and LM-separable, therefore guaranteeing recoverability by Theorem 4.2.




Figure 1: A clustering with 500 points that is weakly separable and LM-separable, but not strongly separable for  $N_p = 5$ , implying that recovery by our algorithm (“MSE”) is guaranteed. We compare against spectral clustering (scikit-learn) with the 4-nearest neighbors affinity matrix because for  $N_p = 5$ , our algorithm uses the 4 nearest neighbors of a point to calculate density.

Our main result states that if  $C$  is weakly separable and LM-separable for a given density parameter  $N_p$ , then it is the unique weakly separable, LM-separable clustering for  $N_p$ , and can be reconstructed efficiently.

**THEOREM 4.2.** *If  $C$  is weakly separable and LM-separable for  $N_p$ , then  $C$  is the unique weakly separable, LM-separable clustering for  $N_p$ , and can be found in  $O(|X|^3 \log(|X|))$  time.*

Since strong separability implies LM-separability and weak separability, expanding an  $\epsilon$ -cut of the dendrogram is not sufficient for finding all weakly separable, LM-separable clusterings by Lemma 4.1.

## 5 Algorithm and Proof of Theorem 4.2

For density parameter  $N_p$ , we will prove that if  $C$  is weakly separable and LM-separable, then Algorithm 1 returns  $C$ , thus proving that  $C$  is the unique weakly separable and LM-separable  $|C|$ -clustering for  $N_p$ . We prove that the algorithm can be implemented to run in  $O(|X|^3 \log(|X|))$  time in the appendix.

Algorithm 2 is an auxiliary algorithm used to define Algorithm 1. Algorithm 2 works by outputting a partial  $K$ -clustering of  $X$  that can then postprocessed to form a  $K$ -clustering. For Algorithm 2,  $A$  governs the maximum variation in density within a partial cluster,  $M$  governs the minimum size of a partial cluster, and  $D$  governs the maximum variation in density among different partial clusters. We use  $C_g(X, N_p, A, M, D)$  to denote the partial clustering of  $X$  that is output by Algorithm 2 with parameters  $N_p, A, M, D$ .

---

**Algorithm 1** Minimal Seed Expansion

---

**Input:**  $X, N_p, M, D, K$ **Output:**  $\hat{C}$ 

- 1:  $C' \leftarrow$  min- $A$  partial clustering of  $X$  with parameters  $N_p, M, D, K$
  - 2:  $\hat{C} \leftarrow$  output of Algorithm 3 with parameters  $X, C', N_p$
- 

At a high level, Algorithm 2 forms a partial clustering of  $X$  by greedily selecting the highest density unclustered point in  $X$ , and creating a maximal cluster centered at the point that satisfies the partial cluster constraints set by the parameters  $A, M, D$ . Note that if there are multiple candidate points for  $x^*$  in a greedy step of Algorithm 2, the final output isn't affected by which candidate is assigned to  $x^*$ . Thus, for parameters  $X, N_p, A, M, D$ , the clustering  $C_g(X, N_p, A, M, D)$  output by Algorithm 2 is unique.

**DEFINITION 5.1.** *The min- $A$  clustering of  $X$  with parameters  $N_p, M, D, K$  is given by  $C_g(X, N_p, A_{\min}, M, D)$  where*

$$A_{\min} = \min_{A \text{ such that } |C_g(X, N_p, A, M, D)|=K} A.$$

*For parameters  $N_p, M, D, K$ , the min- $A$  clustering is denoted by  $C_m(X, N_p, M, D, K)$ .  $A_{\min} \geq 1$  since  $c^*(x, \epsilon) = \{\}$  if  $\epsilon < \epsilon_{N_p}(x)$ .*

Algorithm 1 accepts  $N_p, M, D$  along with a number of clusters  $K$  as input, and begins by finding the min- $A$  partial clustering with parameters  $M, D, K$ . The min- $A$  partial clustering is then passed to Algorithm 3 to produce the final output clustering.

**LEMMA 5.1.** *For a given  $N_p$ , if  $C$  is LM-separable, then  $|C_g(X, N_p, A^\ell(C), 1, \infty)| = K$ .*

*Proof.* We will prove that all partial clusters output by Algorithm 2 are subsets of distinct clusters in  $C$ . Consider the  $i$ th partial cluster output by Algorithm 2. The highest density point  $x$  that the algorithm uses to build the  $i$ th cluster is clearly a local maximum, and can either be from a previously partially reconstructed cluster in  $C$ , or can be from a new cluster in  $C$ . If  $x$  is from a previously partially reconstructed cluster  $c' \in C$ , then it must include some  $z \in X_{c'}^*$  by the definitions of LM-separability and  $A^\ell(C)$ , and the fact that  $\epsilon_{N_p}(x) \geq \epsilon_{N_p}(z)$ . This is a contradiction. This along with the fact that  $M = 1 \leq N_p$  guarantees that  $x$  must be from a new cluster  $c'' \in C$  and  $x \in X_{c''}^*$ . The  $i$ th partial cluster  $c^*(x, A^\ell(C) \cdot \epsilon_{N_p}(x))$  must only include points from  $c''$  by the definition of LM-separated

and  $A^\ell(C)$ . Because only a new cluster center can be added to the partial clustering in each iteration of the algorithm, a partial  $K$ -clustering is output.  $\square$

By Lemma 5.1,  $A_{\min}$  is no larger than  $A^\ell(C)$ .

**LEMMA 5.2.** *For a given  $N_p$ , if  $C$  is weakly separable and LM-separable, the min- $A$  clustering  $C_m(X, N_p, 1, \infty, |C|)$  is extendable to  $C$ .*

*Proof.* Consider the  $i$ th partial cluster output by Algorithm 2. The highest density point  $x$  used by algorithm to build the  $i$ th cluster must be a local maximum, and can either be from a previously partially reconstructed cluster in  $C$ , or can be from a new cluster in  $C$ .

If  $x$  is from a previously partially reconstructed cluster  $c \in C$ , the  $i$ th partial cluster doesn't include any points from other true clusters that are not yet part of a previously output partial cluster. This is because if it did include a point from such a cluster  $c' \in C$ , it would include a point  $z \in X_{c'}^*$  by weak separability of  $C$ , and since  $\epsilon_{N_p}(z) \geq \epsilon_{N_p}(x)$ , it would imply that  $A_{\min} \cdot \epsilon_{N_p}(z) \geq \epsilon(c', c)$  which violates LM-separability of  $C$  since  $A_{\min} \leq A^\ell(C)$ .

If the highest density point  $x$  is from a new true cluster  $c \in C$ , by LM-separability of  $C$ , only points from  $c$  are added to the partial cluster since  $A_{\min} \leq A^\ell(C)$ .

Therefore, the algorithm outputs non-overlapping partial clusters such that each partial cluster chosen doesn't include any points from true clusters that are not yet part of a previously output partial cluster. Observe that, if no points in a cluster  $c$  have been added to previously output partial clusters by the time  $w \in X_c^*$  is selected as the highest density unclustered point by the algorithm, then a partial cluster centered at  $w$  that is a non-empty subset of  $c$  will be output by the algorithm in that step since  $M = 1 \leq N_p$ ,  $A_{\min} \leq A^\ell(C)$  and  $C$  is LM-separable. Therefore, we obtain a partial clustering that contains a partial cluster centered at  $w \in X_c^*$  that is a non-empty subset of  $c$  for each  $c \in C$ . This is a min- $A$  partial clustering for  $K$ , so the algorithm will output a partial  $K$ -clustering that contains a partial cluster centered at some  $w \in X_c^*$  that is a non-empty subset of  $c$  for each  $c \in C$ . Thus, the algorithm outputs a partial  $K$ -clustering that is extendable to  $C$ .  $\square$

**LEMMA 5.3.** *For a given  $N_p$ , if  $C$  is weakly separable, and Algorithm 3 is initialized with a partial clustering  $\hat{C}$  that is extendable to  $C$ , then Algorithm 3 recovers  $C$ .*

*Proof.* Suppose that at some step in Algorithm 3, a point  $x$  is added to a cluster  $c' \neq c(x)$  with some  $\epsilon$ . This implies that at this step, there is no point  $y \in c(x)$

that is not yet in  $\hat{C}$  that is  $\epsilon$ -connected to another point  $z \in c(x)$  that is already in  $\hat{C}$ . This is a contradiction to the weakly separated assumption on  $C$ .  $\square$

Lemma 5.2 and Lemma 5.3, imply if  $C$  is weakly separable and LM-separable, then it is the unique weakly separable, LM-separable clustering for  $X$  and  $N_p$ . As a consequence of Lemmas 7.1 and 7.2 in the appendix, Algorithm 1 can be implemented in  $O(|X|^3 \log(|X|))$  time.

---

**Algorithm 2** Greedy algorithm to find partial clusters

---

**Input:**  $X, N_p, A, M, D$

**Output:**  $C'$

```

 $C' \leftarrow \emptyset$ 
MinExtracted  $\leftarrow \infty$ 
Tried  $\leftarrow \emptyset$ 
while  $X \setminus \text{Tried} \neq \emptyset$  do
   $x^* \leftarrow \arg \min_{x \in X \setminus \text{Tried}} \epsilon_{N_p}(x)$ 
  if  $\epsilon_{N_p}(x^*) > D \cdot \text{MinExtracted}$  then
    break loop
  end if
   $c' \leftarrow c^*(x^*, A \cdot \epsilon_{N_p}(x^*))$ 
  if  $|c'| \geq M$  and  $c' \cap c = \emptyset \ \forall c \in C'$  then
    add  $c'$  to  $C'$ 
    if MinExtracted =  $\infty$  then
      MinExtracted  $\leftarrow \epsilon_{N_p}(x^*)$ 
    end if
     $X \leftarrow X \setminus c'$ 
  else
    Tried  $\leftarrow \text{Tried} \cup x$ 
  end if
end while

```

---



---

**Algorithm 3** Greedy algorithm to expand partial clusters

---

**Input:**  $X, C'$

**Output:**  $\hat{C}$

```

 $\hat{C} \leftarrow C'$ 
 $Y \leftarrow X \setminus (\cup_{c \in C'} c)$ 
while  $\cup_{c \in \hat{C}} c \neq X$  do
   $(x^*, c^*) \leftarrow \arg \min_{x \in Y, c \in \hat{C}} \epsilon(x, c)$ 
   $c^* \leftarrow c^* \cup \{x^*\}$ 
   $Y \leftarrow Y \setminus x^*$ 
end while

```

---

## 6 Experiments

We compare an implementation of Algorithm 1 to state of the art algorithms for non-convex cluster recovery on a range of datasets. In this implementation, we use a slight variation of Algorithm 2 where each time a cluster

is output, the points in the cluster are removed from  $X$ . This modification makes the check of whether a new cluster intersects with a previously output cluster inapplicable since only unclustered points remain in  $X$  at the beginning of each greedy step. We use this approach because it appears to work well on datasets that have substantial cluster overlap. To increase speed, instead of finding the min- $A$  clustering exactly, the implemented algorithm approximates the min- $A$  clustering by progressively adjusting  $A$  until a  $K$ -clustering is output by the modified version of Algorithm 2. To improve performance slightly, we use  $N_p = 2$  in Algorithm 3, regardless of the  $N_p$  used for finding the initial partial clusters.

We compare our algorithm to  $K$ -means, spectral clustering, HDBSCAN, and OPTICS. We refer to our algorithm as “MSE” which stands for minimal seed expansion, and implemented it in Python. We use the spectral clustering and  $K$ -means implementations from the scikit-learn, where the number of clusters  $K$  is specified by the user. The affinity matrix in spectral clustering is formed using each point’s  $K_n$  nearest neighbors, where  $K_n$  is by the user. The implementation of  $K$ -means uses the “greedy  $K$ -means++” algorithm [17, 18]. We use the HDBSCAN implementation from the Python HDBSCAN clustering library, and we run it using the default cluster selection criteria named Excess of Mass (eom). The default setting for HDBSCAN sets the  $N_p$  equal to  $M+1$  where  $M$  is the minimum possible cluster size set by the user. We refer to this default version as “HDBSCAN (d).” We refer to the version of HDBSCAN where we set  $N_p$  and  $M$  independently as “HDBSCAN.” We use the OPTICS implementation from scikit-learn, where  $N_p$ , the minimum cluster size  $M$ , and the cluster selection parameter  $X_i$  are set by the user. Neither HDBSCAN or OPTICS uses the number of clusters  $K$ .

The Adjusted Rand Index (ARI) and Normalized Mutual Information (NMI) are the most common measures of similarity between an estimated clustering and the true clustering, and we report both for each dataset tested. The algorithms’ performance is compared on the real world benchmark datasets whose properties are in Table 2 in the appendix. All of these datasets are available on the UCI server [19]. “Cancer” refers to the Breast Cancer Wisconsin (Diagnostic) dataset. “Digits” refers to the test set of the Optical Recognition of Handwritten Digits dataset. “Letters” refers to the test set of the Letter Recognition dataset. “MNIST” refers to the test set of the MNIST dataset. For the MNIST dataset, we used t-SNE to reduce the dimensionality to two [20]. We also compare the same algorithms for 5 artificial datasets which were used in the example titled “Comparing different clustering algorithms on toy

datasets” on the scikit-learn website [21].

The results for benchmark datasets are reported in Table 1, and results for artificial datasets are in Table 3 in the appendix. For our algorithm, we did not optimize  $N_p$ ,  $M$  and  $D$  in a meaningful way. For all experiments we set  $N_p = 3$ , picked  $M$  values that seemed reasonable, and used the  $D$  value from the set  $\{1.5, 2, 10\}$  that gave the best results. For spectral clustering, we tried all  $K_n$  in the range  $\{1, 2, \dots, 20\}$  and reported the clustering with the best ARI. For HDBSCAN (d), we tried all  $N_p$  in the range  $\{2, 3, \dots, 20\}$  and reported the clustering with best ARI. For HDBSCAN and OPTICS, we used the same  $M$  and  $N_p$  we used for our algorithm. For OPTICS, we tried all  $X_i$  values in the set  $\{0, 0.05, 0.1, \dots, 0.95\}$  and reported the clustering with best ARI.

For spectral clustering, in order of dataset appearance in the Table 1, the chosen  $K_n$  values are 4, 6, 18, 3, 5, 4, 19, 11. For HDBSCAN (d), in order of appearance in the Table 1, the chosen  $N_p$  values are 3, 20, 6, 3, 6, 4, 3, 14. For our algorithm, HDBSCAN, and OPTICS,  $N_p$  was set to 3 for all datasets and  $M$  was set to 35 for Iris, 30 for Wine, 50 for Seeds, 8 for Glass, 100 for Cancer and Digits, and 70 for Letters. For OPTICS, in order of appearance in the Table 1, the chosen  $X_i$  values are 0, 0.25, 0, 0.2, 0.05, 0, 0, 0. For our algorithm,  $D$  was set to 20 for Iris, Wine, Glass, and Cancer, 2 for Seeds, Letters and MNIST, and 1.5 for Digits.

Observe that with just loose lower bound on the minimum cluster size ( $M$ ), and some knowledge of the maximum difference in maximum density among different clusters ( $D$ ), our algorithm gives as good or better ARI than the other algorithms on all artificial and benchmark datasets except the Glass dataset. On the Glass dataset, HDBSCAN (d) gives the best ARI, but outputs an incorrect number of clusters. The NMI values obtained by our algorithm are also competitive. As an example, in the case of the Digits dataset, it is clear in Figure 2 in the appendix that the our algorithm does a significantly better job of recovering the Digits clustering than the other algorithms. While our implementation of MSE was not optimized for speed, MSE clustered each dataset in roughly three minutes or less on a laptop computer with 32 threads and 64 GB of RAM.

For our algorithm,  $D$  (and similarly  $M$  and  $N_p$ ) can be optimized by running our algorithm for various  $D$  values and using the one that gives the clustering with the best DBCV score [22] (or other cluster quality score), though we leave this optimization for future work.

| dataset | algorithm   | ARI   | NMI   | $ \hat{C} $ |
|---------|-------------|-------|-------|-------------|
| Iris    | MSE         | 0.886 | 0.871 | 3           |
|         | Spectral    | 0.835 | 0.833 | 3           |
|         | $K$ -means  | 0.716 | 0.742 | 3           |
|         | HDBSCAN (d) | 0.568 | 0.734 | 2           |
|         | HDBSCAN     | 0.568 | 0.734 | 2           |
|         | OPTICS      | 0.732 | 0.753 | 4           |
| Wine    | MSE         | 0.439 | 0.430 | 3           |
|         | Spectral    | 0.401 | 0.395 | 3           |
|         | $K$ -means  | 0.371 | 0.429 | 3           |
|         | HDBSCAN (d) | 0.292 | 0.379 | 2           |
|         | HDBSCAN     | 0.291 | 0.403 | 3           |
|         | OPTICS      | 0.327 | 0.428 | 2           |
| Seeds   | MSE         | 0.725 | 0.682 | 3           |
|         | Spectral    | 0.657 | 0.660 | 3           |
|         | $K$ -means  | 0.717 | 0.695 | 3           |
|         | HDBSCAN (d) | 0.336 | 0.468 | 5           |
|         | HDBSCAN     | 0.409 | 0.469 | 3           |
|         | OPTICS      | 0.547 | 0.594 | 3           |
| Glass   | MSE         | 0.191 | 0.374 | 7           |
|         | Spectral    | 0.202 | 0.367 | 7           |
|         | $K$ -means  | 0.216 | 0.388 | 7           |
|         | HDBSCAN (d) | 0.277 | 0.446 | 6           |
|         | HDBSCAN     | 0.222 | 0.395 | 6           |
|         | OPTICS      | 0.264 | 0.440 | 4           |
| Cancer  | MSE         | 0.743 | 0.628 | 2           |
|         | Spectral    | 0.583 | 0.487 | 2           |
|         | $K$ -means  | 0.491 | 0.465 | 2           |
|         | HDBSCAN (d) | 0.625 | 0.487 | 4           |
|         | HDBSCAN     | 0.000 | 0.000 | 1           |
|         | OPTICS      | 0.737 | 0.620 | 2           |
| Digits  | MSE         | 0.864 | 0.898 | 10          |
|         | Spectral    | 0.781 | 0.892 | 10          |
|         | $K$ -means  | 0.615 | 0.731 | 10          |
|         | HDBSCAN (d) | 0.575 | 0.770 | 22          |
|         | HDBSCAN     | 0.559 | 0.762 | 9           |
|         | OPTICS      | 0.259 | 0.628 | 10          |
| Letters | MSE         | 0.193 | 0.447 | 26          |
|         | Spectral    | 0.098 | 0.408 | 26          |
|         | $K$ -means  | 0.130 | 0.356 | 26          |
|         | HDBSCAN (d) | 0.023 | 0.536 | 463         |
|         | HDBSCAN     | 0.003 | 0.064 | 3           |
|         | OPTICS      | 0.047 | 0.343 | 24          |
| MNIST   | MSE         | 0.854 | 0.854 | 10          |
|         | Spectral    | 0.632 | 0.746 | 10          |
|         | $K$ -means  | 0.656 | 0.751 | 10          |
|         | HDBSCAN (d) | 0.712 | 0.795 | 11          |
|         | HDBSCAN     | 0.603 | 0.766 | 8           |
|         | OPTICS      | 0.444 | 0.689 | 11          |

Table 1:  $|\hat{C}|$  is the number of clusters output by the algorithm. For HDBSCAN and OPTICS, the set of noise points counts as a cluster.



## References

- [1] J. Hou, H. Gao, and X. Li, "Dsets-dbscan: A parameter-free clustering algorithm," *IEEE Transactions on Image Processing*, vol. 25, no. 7, pp. 3182–3193, 2016.
- [2] D. Birant and A. Kut, "St-dbscan: An algorithm for clustering spatial-temporal data," *Data & knowledge engineering*, vol. 60, no. 1, pp. 208–221, 2007.
- [3] M. Stonebraker, J. Frew, K. Gardels, and J. Meredith, "The sequoia 2000 storage benchmark," *ACM SIGMOD Record*, vol. 22, no. 2, pp. 2–11, 1993.
- [4] R. Ding, Q. Wang, Y. Dang, Q. Fu, H. Zhang, and D. Zhang, "Yading: Fast clustering of large-scale time series data," *Proceedings of the VLDB Endowment*, vol. 8, no. 5, pp. 473–484, 2015.
- [5] U. Von Luxburg, "A tutorial on spectral clustering," *Statistics and computing*, vol. 17, pp. 395–416, 2007.
- [6] M. Ester, H.-P. Kriegel, J. Sander, X. Xu, *et al.*, "A density-based algorithm for discovering clusters in large spatial databases with noise," in *kdd*, vol. 96, pp. 226–231, 1996.
- [7] M. Ankerst, M. M. Breunig, H.-P. Kriegel, and J. Sander, "Optics: Ordering points to identify the clustering structure," *ACM Sigmod record*, vol. 28, no. 2, pp. 49–60, 1999.
- [8] L. McInnes, J. Healy, S. Astels, *et al.*, "hdbscan: Hierarchical density based clustering.," *J. Open Source Softw.*, vol. 2, no. 11, p. 205, 2017.
- [9] P. Bhattacharjee and P. Mitra, "A survey of density based clustering algorithms," *Frontiers of Computer Science*, vol. 15, pp. 1–27, 2021.
- [10] Y. Zhu, K. M. Ting, and M. J. Carman, "Density-ratio based clustering for discovering clusters with varying densities," *Pattern Recognition*, vol. 60, pp. 983–997, 2016.
- [11] A. Ram, S. Jalal, A. S. Jalal, and M. Kumar, "A density based algorithm for discovering density varied clusters in large spatial databases," *International Journal of Computer Applications*, vol. 3, no. 6, pp. 1–4, 2010.
- [12] Z. Wang, Z. Ye, Y. Du, Y. Mao, Y. Liu, Z. Wu, and J. Wang, "Amd-dbscan: An adaptive multi-density dbscan for datasets of extremely variable density," in *2022 IEEE 9th International Conference on Data Science and Advanced Analytics (DSAA)*, pp. 1–10, IEEE, 2022.
- [13] S. Hess, W. Duivesteijn, P. Honysz, and K. Morik, "The spectacl of nonconvex clustering: A spectral approach to density-based clustering," in *Proceedings of the AAAI conference on artificial intelligence*, vol. 33, pp. 3788–3795, 2019.
- [14] J. Jang and H. Jiang, "Dbscan++: Towards fast and scalable density clustering," in *International conference on machine learning*, pp. 3019–3029, PMLR, 2019.
- [15] H. Jiang, "Density level set estimation on manifolds with dbscan," in *International Conference on Machine Learning*, pp. 1684–1693, PMLR, 2017.
- [16] H. Esfandiari, V. Mirrokni, and P. Zhong, "Almost linear time density level set estimation via dbscan," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, pp. 7349–7357, 2021.
- [17] R. Ostrovsky, Y. Rabani, L. J. Schulman, and C. Swamy, "The effectiveness of lloyd-type methods for the k-means problem," *Journal of the ACM (JACM)*, vol. 59, no. 6, pp. 1–22, 2013.
- [18] D. Arthur and S. Vassilvitskii, "k-means++: the advantages of careful seeding," *SODA '07, (USA)*, p. 1027–1035, Society for Industrial and Applied Mathematics, 2007.
- [19] D. Dua and C. Graff, "Uci machine learning repository," 2017.
- [20] L. Van der Maaten and G. Hinton, "Visualizing data using t-sne.," *Journal of machine learning research*, vol. 9, no. 11, 2008.
- [21] "Comparing different clustering algorithms on toy datasets." [https://scikit-learn.org/dev/autodocs/examples/cluster/plot\\_cluster\\_comparison.html#sphx-glr-auto-examples-cluster-plot-cluster-comparison-py](https://scikit-learn.org/dev/autodocs/examples/cluster/plot_cluster_comparison.html#sphx-glr-auto-examples-cluster-plot-cluster-comparison-py). Accessed: 2024-10-03.
- [22] D. Moulavi, P. A. Jaskowiak, R. J. G. B. Campello, A. Zimek, and J. Sander, "Density-based clustering validation," in *Proceedings of the 2014 SIAM International Conference on Data Mining, Philadelphia, Pennsylvania, USA, April 24-26, 2014* (M. J. Zaki, Z. Obradovic, P. Tan, A. Banerjee, C. Kamath, and S. Parthasarathy, eds.), pp. 839–847, SIAM, 2014.

## 7 Appendix

### 7.1 Proofs

LEMMA 7.1. *For a given  $X, N_p, M, D$ , the min- $A$  clustering can be found in  $O(|X|^3 \log(|X|))$  time.*

*Proof.* We implement the greedy step of Algorithm 2 to have a deterministic rule for deciding whether to choose  $x^*$  to be  $x$  or  $y$  if  $\epsilon_{N_p}(x) = \epsilon_{N_p}(y)$  (such as taking the point that comes first in the dataset).

For a given  $X, N_p, M, D$ , the number of clusters in  $C_g(X, N_p, A, M, D)$  monotonically decreases as  $A$  increases because of the following property. Consider some  $a, a' \geq 1$  such that  $a > a'$ . At the end of the  $i$ th iteration of Algorithm 2, the set of remaining candidates for future  $x^*$  where  $c^*(x^*, A \cdot \epsilon_{N_p}(x^*))$  will not intersection with a previously output cluster in the case when  $A = a'$  is a superset of the set of such candidates when  $A = a$ .

We will prove this by induction. Suppose this property holds for the  $(i-1)$ th iteration of Algorithm 2, and consider the  $i$ th iteration. The point  $x_{a'}$  picked to be  $x^*$  by the algorithm if  $A = a'$  may or may not already be included in a cluster previously output by the algorithm with  $A = a$ . If not, then  $x_{a'}$  will be chosen as  $x^*$  for the algorithm with  $A = a$  by the inductive hypothesis, so clearly at the end of the  $i$ th step, the property still holds since  $a > a'$ .

If on the other hand,  $x_{a'}$  is included in a cluster previously output by the algorithm with  $A = a$ , then all points  $x$  that are  $(A \cdot \epsilon_{N_p}(x))$ -connected to  $x_{a'}$  already cannot be candidates at the beginning of the  $i$ th step for the algorithm with  $A = a$ . Therefore, in this case, at the end of the  $i$ th step, the property still holds.

Due to the monotonicity property we have just proved, if we have a set  $S$  that includes all  $A$  values that lead to distinct clusterings  $C_g(X, N_p, A, M, D)$ , then we can use binary search on  $S$  to find the min- $A$  clustering for  $X, N_p, M, D$ . Algorithm 2 runs in  $O(|X|^2)$  time. Therefore, if we have such a set  $S$  and the sorted list of its elements, this binary search approach to find the min- $A$  clustering runs in  $O(\log(|S|) \cdot |X|^2)$  time. Sorting the elements in  $S$  for use in binary search runs in  $O(|S| \log(|S|))$  time.

Consider a point  $x^*$  selected in a greedy step of Algorithm 2. One set that includes all  $A$  values that could lead to different clusterings is given by  $T_{x^*} = \{d(y, z)/\epsilon_{N_p}(x^*) : y, z \in X\}$  i.e. the set of all distances between points in  $X$  divided by  $\epsilon_{N_p}(x^*)$ . Thus,  $S = \cup_{x \in X} T_x$  includes every possible  $A$  value that could lead to a different clustering. We have that  $|S| = O(|X|^3)$ . Thus, the binary search to find the min- $A$  clustering runs in  $O(|X|^2 \log(|X|))$  time. Constructing  $S$  runs in  $O(|X|^3)$  time, and sorting the values of  $S$  for binary

| dataset | # points | # features | # clusters |
|---------|----------|------------|------------|
| Iris    | 150      | 4          | 3          |
| Wine    | 178      | 13         | 3          |
| Seeds   | 210      | 7          | 3          |
| Glass   | 214      | 9          | 7          |
| Cancer  | 556      | 30         | 2          |
| Digits  | 1,797    | 64         | 10         |
| Letter  | 4,000    | 16         | 26         |
| MNIST   | 10,000   | 784        | 10         |

Table 2: benchmark datasets and their properties

search runs in  $O(|X|^3 \log(|X|))$  time. The total runtime of the approach is therefore  $O(|X|^3 \log(|X|))$ .  $\square$

LEMMA 7.2. *For a given  $X, N_p, \hat{C}$ , Algorithm 3 can be implemented to run in  $O(|X|^2)$  time.*

*Proof.* Observe that for each greedy step of Algorithm 3, any point  $x \in X \setminus (\cup_{c \in \hat{C}} c)$  and cluster  $c \in \hat{C}$  that minimize the quantity  $\epsilon^!(x, c) = \min_{y \in c} \max(d(x, y), \epsilon_{N_p}(x), \epsilon_{N_p}(y))$  can be selected as  $(x^*, c^*)$ . To see this, consider the set  $S$  of unclustered points that are closest to some cluster in terms of  $\epsilon$ -distance. Denote this minimum  $\epsilon$ -distance by  $\epsilon'$ . Clearly,  $S$  must include an unclustered point that contains a clustered point in its  $\epsilon'$ -ball. Thus, at each greedy step of Algorithm 3, we can pick  $(x, c)$  that minimizes  $\epsilon^!(x, y)$ .

To initialize the algorithm, for each unclustered point  $x \in X \setminus (\cup_{c \in \hat{C}} c)$ , and each clustered point  $y$ , we compute  $\max(d(x, y), \epsilon_{N_p}(x), \epsilon_{N_p}(y))$  which then allows us to compute  $\epsilon^!(x) = \min_{c \in \hat{C}} \epsilon^!(x, c)$  and  $c^!(x) = \arg \min_{c \in \hat{C}} \epsilon^!(x, c)$  for every  $x \in X \setminus (\cup_{c \in \hat{C}} c)$ . This can be done in  $O(|X|^2)$  time.

For each greedy step of Algorithm 3, we simply set  $(x^*, c^*)$  equal to the  $(x, c)$  that minimizes  $\arg \min_{c \in \hat{C}} \epsilon^!(x, c)$  in  $O(|X|)$  time by checking  $(\epsilon^!(x), c^!(x))$  for all  $x \in X \setminus (\cup_{c \in \hat{C}} c)$ . After assigning  $x^*$  to  $c^*$ , we update  $\epsilon^!(x)$  for every remaining unclustered point  $x$  by setting  $\epsilon^!(x) \leftarrow \min(\epsilon^!(x), \max(d(x, x^*), \epsilon_{N_p}(x), \epsilon_{N_p}(x^*)))$ , and setting  $c^!(x) \leftarrow c^*$  if the value of  $\epsilon^!(x)$  is changed. Each of these greedy steps takes  $O(|X|)$  time and there are at most  $O(|X|)$  greedy steps. This stage of the algorithm therefore runs in  $O(|X|^2)$  time.  $\square$

**7.2 Experiments** For spectral clustering, in order of appearance in Table 3, the chosen  $K_n$  values are 5, 5, 18, 6, 15. For HDBSCAN (d), in order of appearance in Table 3, the chosen  $N_p$  values are 3, 6, 9, 9, 14. For our algorithm, HDBSCAN, and OPTICS,  $N_p$  was

| dataset               | algorithm       | ARI    | NMI   | $ \hat{C} $ |
|-----------------------|-----------------|--------|-------|-------------|
| Two Circles           | MSE             | 1.000  | 1.000 | 2           |
|                       | Spectral        | 1.000  | 1.000 | 2           |
|                       | <i>K</i> -means | -0.002 | 0.000 | 2           |
|                       | HDBSCAN (d)     | 1.000  | 1.000 | 2           |
|                       | HDBSCAN         | 1.000  | 1.000 | 2           |
|                       | OPTICS          | 1.000  | 1.000 | 2           |
| Two Moons             | MSE             | 1.000  | 1.000 | 2           |
|                       | Spectral        | 1.000  | 1.000 | 2           |
|                       | <i>K</i> -means | 0.233  | 0.176 | 2           |
|                       | HDBSCAN (d)     | 1.000  | 1.000 | 2           |
|                       | HDBSCAN         | 0.699  | 0.737 | 4           |
|                       | OPTICS          | 1.000  | 1.000 | 2           |
| Fixed Variance Blobs  | MSE             | 0.964  | 0.948 | 3           |
|                       | Spectral        | 0.976  | 0.961 | 3           |
|                       | <i>K</i> -means | 0.970  | 0.954 | 3           |
|                       | HDBSCAN (d)     | 0.867  | 0.847 | 4           |
|                       | HDBSCAN         | 0.568  | 0.729 | 3           |
|                       | OPTICS          | 0.569  | 0.733 | 2           |
| Anisotropic           | MSE             | 1.000  | 1.000 | 3           |
|                       | Spectral        | 0.994  | 0.989 | 3           |
|                       | <i>K</i> -means | 0.555  | 0.593 | 3           |
|                       | HDBSCAN (d)     | 0.917  | 0.888 | 4           |
|                       | HDBSCAN         | 0.991  | 0.983 | 4           |
|                       | OPTICS          | 0.970  | 0.951 | 4           |
| Varied Variance Blobs | MSE             | 0.896  | 0.867 | 3           |
|                       | Spectral        | 0.896  | 0.873 | 3           |
|                       | <i>K</i> -means | 0.787  | 0.778 | 3           |
|                       | HDBSCAN (d)     | 0.808  | 0.809 | 4           |
|                       | HDBSCAN         | 0.844  | 0.817 | 4           |
|                       | OPTICS          | 0.906  | 0.865 | 3           |

Table 3:  $|\hat{C}|$  is the number of clusters output by the algorithm. For HDBSCAN and OPTICS, the set of noise points counts as a cluster.

set to 3 for all datasets and  $M$  was set to 60 for all datasets in Table 3. For OPTICS, in order of appearance in the Table 3, the chosen  $X_i$  values are 0.15, 0.1, 0.3, 0.2, 0.15. For our algorithm,  $D$  was set to 10 for Varied Variance Blobs, and 2 for all other datasets in Table 3.

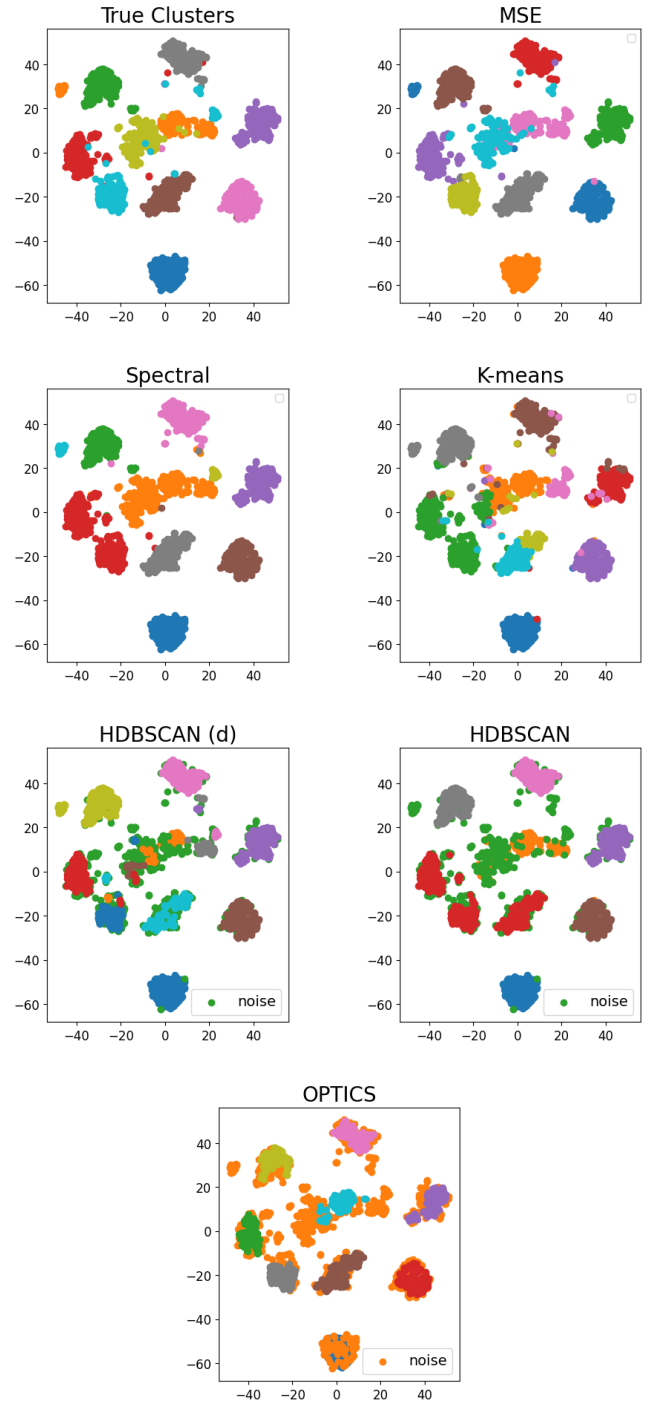


Figure 2: t-SNE plots for the Digits dataset.

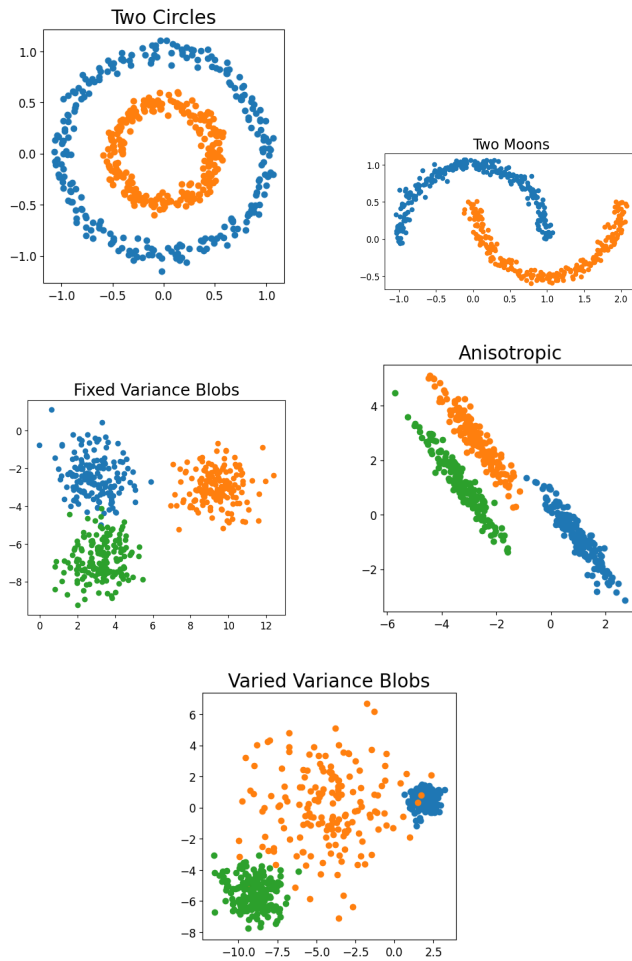


Figure 3: Artificial datasets used in Table 3 .